

# Sovereign LLM Operations with RAGIX/KOAS

## A Framework for Confidential Document Analysis in Regulated Environments

**Author:** Olivier Vitrac, PhD, HDR | Adservio Innovation Lab **Version:** 1.2.0 **Date:** 2026-01-29  
**Classification:** Technical Architecture Document

### Abstract

This document presents the architectural principles and operational procedures enabling Large Language Models (LLMs) to perform document analysis tasks in environments where data confidentiality is paramount. The RAGIX (*Retrieval-Augmented Generative Interactive eXecution*) framework, combined with the KOAS (*Kernel-Orchestrated Audit System*) methodology, provides a sovereign alternative to cloud-based AI services. We demonstrate how policy enforcement, deterministic kernels, and local LLM deployment create an auditable system suitable for sensitive document processing—including technical audits, compliance verification, and specification analysis—without exposing data to external services.

**Keywords:** Sovereign AI, Local LLM, RAG, Document Analysis, Data Sovereignty, Air-Gapped Systems, Compliance, KOAS

### Pipeline Invariants (MUST / MUST NOT)

This section defines non-negotiable rules for any LLM or agent operating within this framework.

#### Caution

#### MUST

1. **Strip Markdown metadata in all external outputs** — YAML/TOML/JSON front-matter, fenced `metadata` blocks, and HTML provenance comments (`<!-- PROVENANCE ... -->`) are internal-only artifacts. External reports receive `content_clean` exclusively.
2. **Enforce output-level contract via CLI** — `--output-level=external` MUST strip metadata, redact paths, anonymize identifiers. This is enforced by build-time validation, not trust.
3. **Forced caching for all document LLM calls** — Every LLM call in `doc_*` kernels MUST be logged with:
  1. `call_hash` = SHA256 of canonical request JSON (sorted keys, normalized whitespace, volatile fields stripped)
  2. `inputs_merkle_root` = Merkle root of ordered child hashes for pyramidal synthesis
  3. Cache key includes: `model_id`, `temperature`, `template_id@version`, `prompt_canonical`
4. **Protect fenced/inline code inside documents** — Boilerplate detection patterns MUST NOT match content within triple-backtick fences or inline code spans. Code blocks are preserved verbatim.
5. **Canonical ordering for reproducibility** — Pyramid children MUST be ordered deterministically (by document path, then chunk index) before Merkle root computation. Non-deterministic ordering invalidates cache.
6. **Seed logging for partial-determinism kernels** — Clustering kernels

(`doc_cluster`, `doc_cluster_leiden`, `partition`) MUST log their random seed in the audit trail. Same seed + same input = same output.

## MUST NOT

1. **Apply boilerplate cleaning to code or Code+Docs mode** — Boilerplate detection is for **pure document corpora only**. In mixed mode, code files bypass boilerplate filtering entirely.
2. **Include provenance markers in external outputs** — Keys like `call_hash`, `inputs_merkle_root`, `run_id`, `endpoint`, `model` are denylist items. Build fails if detected in external report.
3. **Modify representative content in `doc_extract`** — Quality filtering is achieved by **penalizing boilerplate patterns in `_score_sentence_quality()`**, not by cleaning/transforming the text. Original content is preserved.
4. **Allow orchestrator to see document excerpts** — `--output-level=orchestrator` returns metrics only: kernel names, timing, success/failure, aggregate counts. No text content.
5. **Trust implicit defaults for isolation** — Isolation settings MUST be explicit in manifest or CLI. Absent flags do not imply "safe"; they imply "internal mode."

## Table of Contents

1. Introduction
2. Sovereignty Architecture
3. The KOAS Paradigm: Kernels Compute, LLMs Interpret
4. KOAS Kernel Inventory and Capabilities
5. Policy Enforcement Mechanisms
6. Operational Modes
7. Gray Environment Deployment
8. Attestation and Audit Trail
9. Practical Implementation
10. Case Study: Metropolitan Infrastructure Audit
11. Limitations and Mitigations
12. Conclusion
13. References

## 1. Introduction

### 1.1 The Confidentiality Challenge

Organizations increasingly require AI-assisted analysis of sensitive documents: technical specifications, audit reports, contractual documents, and compliance assessments. Traditional cloud-based LLM services (OpenAI, Anthropic API, Google AI) present fundamental confidentiality concerns:

Concern	Cloud LLM Risk	Impact
Data exfiltration	Documents sent to external servers	Confidentiality breach
Training data leakage	User data may train future models	IP loss
Jurisdiction	Data processed in foreign territories	Regulatory violation
Audit opacity	No visibility into processing	Compliance failure

## 1.2 The Sovereignty Imperative

For regulated industries (finance, healthcare, government, critical infrastructure), data sovereignty is not optional. The European GDPR, French ANSSI guidelines, and sector-specific regulations mandate:

1. **Data localization** — Processing must occur within controlled perimeters
2. **Auditability** — Complete traceability of data transformations
3. **Minimal exposure** — Data access limited to necessary operations
4. **Provenance tracking** — Origin and transformation history preserved

## 1.3 RAGIX/KOAS Solution

This document demonstrates how the RAGIX/KOAS stack achieves these requirements through:

- **Local LLM deployment** via Ollama (no external API calls)
- **Deterministic kernel computation** (reproducible, auditable)
- **Policy-enforced sandboxing** (controlled execution environment)
- **Cryptographic audit trails** (tamper-evident logging)
- **Cache-based replay** (LLM-free verification possible)

## 2. Sovereignty Architecture

### 2.1 System Topology

The sovereign architecture ensures all data processing occurs within a controlled boundary:

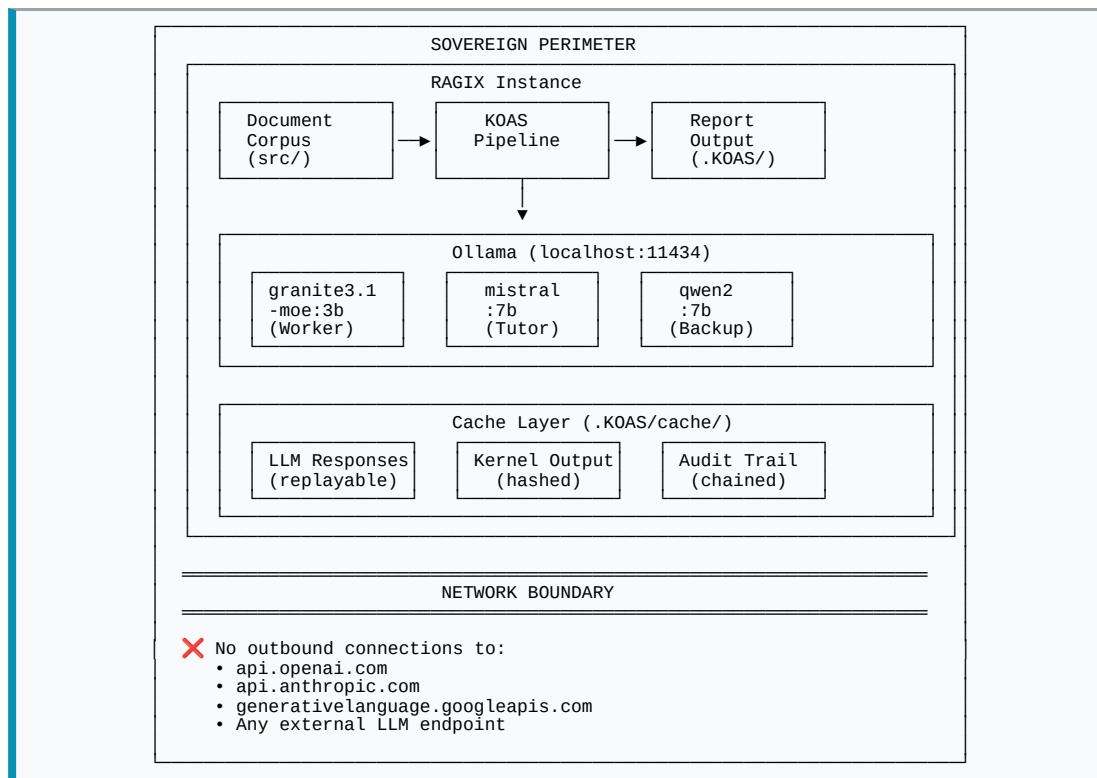


Figure 1. Isolation of data/documents/codes and of their processing.

## 2.2 Core-Shell (Onion) Architecture with External Orchestrator

For maximum isolation, RAGIX supports a **core-shell architecture** where an external orchestrating LLM (e.g., Claude, GPT-4) can drive the analysis without ever accessing the raw data. This "onion" model provides layered isolation.

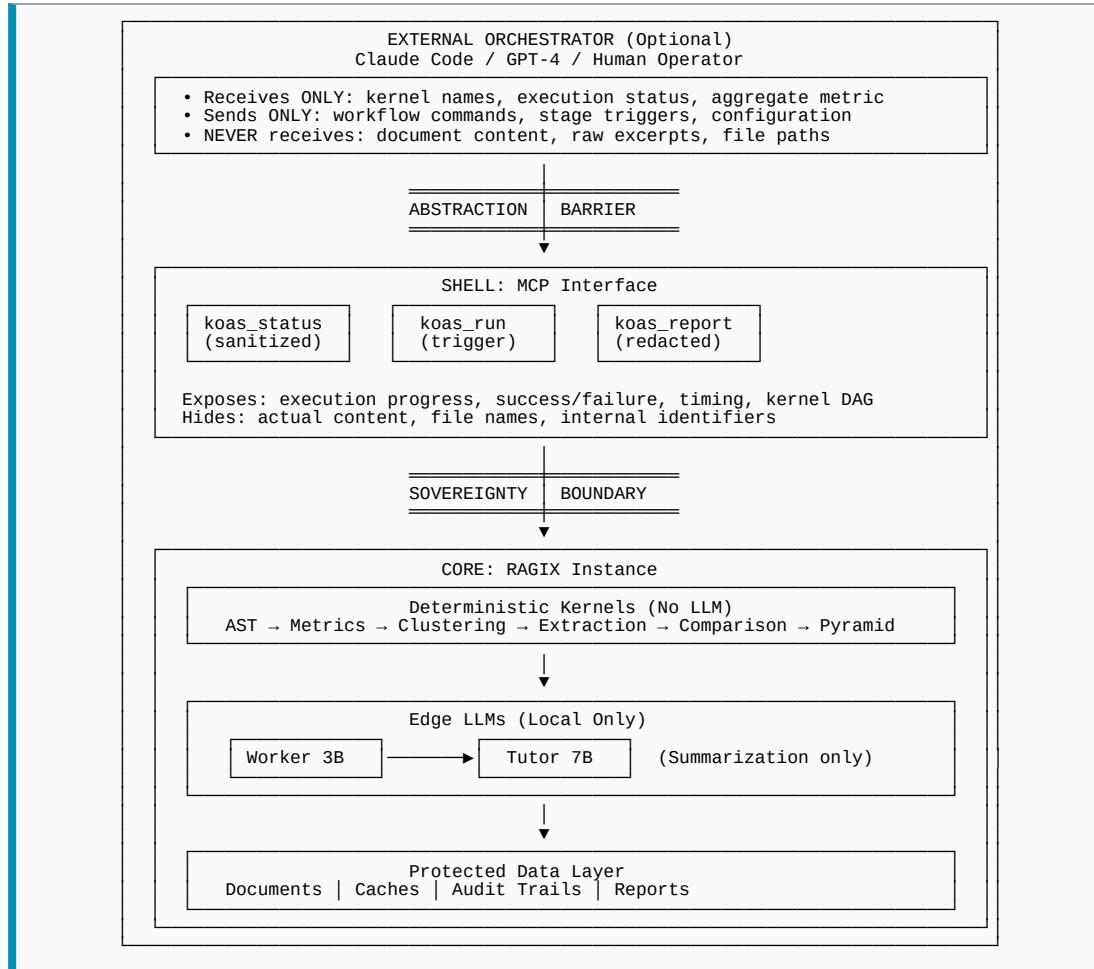


Figure 2. Core-Shell Architecture with External Orchestrator

### Isolation Guarantees of the Onion Model:

Layer	Sees	Does Not See	Control Level
External Orchestrator	Kernel names, metrics, timing	Content, file names, excerpts	Commands only
MCP Shell	Sanitized status, redacted reports	Raw paths, internal IDs	Trigger execution
RAGIX Core	Full document content	External systems	Full processing
Edge LLMs	Pre-structured prompts	Network, filesystem	Inference only

This architecture enables scenarios where:

- A cloud-based LLM orchestrates the workflow without data exposure
- Human operators receive progress updates without sensitive content
- Compliance officers audit the process through sanitized logs
- The core remains fully sovereign regardless of the orchestration layer

## 2.3 Data Flow Guarantees

**Theorem (Data Containment):** Under correct configuration, no document content traverses the sovereign perimeter boundary.

**Proof sketch:**

1. All LLM calls are routed exclusively to `localhost:11434` (Ollama)
2. Ollama operates in inference-only mode (no telemetry, no model updates)
3. Kernel computations are pure functions on local filesystem data
4. Cache reads/writes occur only within `.KOAS/` directory
5. No network sockets are opened except to localhost

This can be verified by:

- Network traffic monitoring (`tcpdump`, `wireshark`)
- Process tracing (`strace -e connect`)
- Firewall rules blocking outbound on ports 80, 443

## 2.4 Edge LLM Strategy: Minimal Compute at the Boundary

The KOAS architecture deliberately places LLMs at the **edge** of the computation graph, not at its center. This design choice is driven by both sovereignty and efficiency considerations.

**Why Edge Placement?**

1. **Minimal attack surface** — LLMs only see pre-structured data, not raw documents
2. **Deterministic core** — 90%+ of computation is hash-verifiable
3. **Replaceable models** — Edge LLMs can be swapped without reprocessing
4. **Reduced token consumption** — Kernels compress information before LLM sees it
5. **Parallel scaling** — Multiple documents processed by kernels, LLM only for synthesis

**Token Efficiency Analysis** (example values for 137-document corpus):

Approach	Tokens per Document	Total for 137 docs	Cost Factor
Full-document LLM	~50,000	6.85M tokens	100x
KOAS (kernel → LLM)	~2,000	274K tokens	4x
KOAS + caching	~500 (cache hits)	68K tokens	1x

**Model Selection for Edge Compute:**

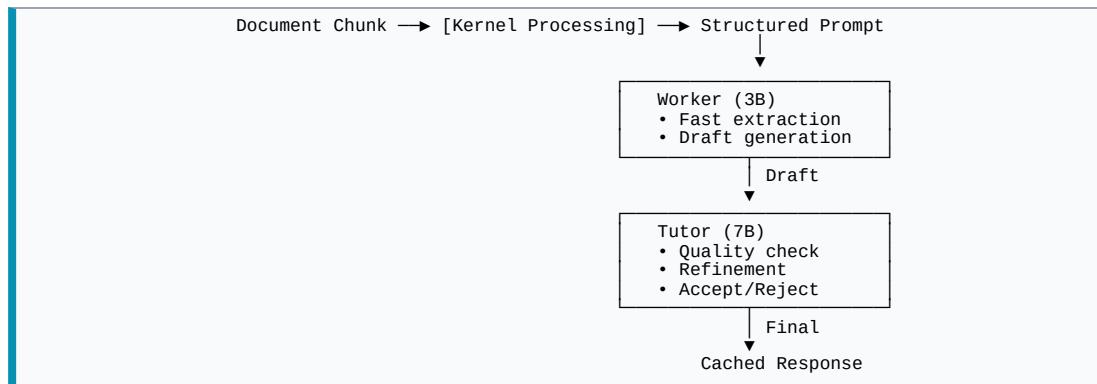
Model	Parameters	VRAM	Role	Justification
granite3.1-moe:3b	3B (MoE)	~4GB	Worker	Sparse activation: only 2B params active per forward pass. Ideal for bulk extraction where speed > quality.
mistral:7b-instruct	7B	~6GB	Tutor	Strong instruction-following. Validates and refines worker outputs. Quality gate.
llama3:8b	8B	~8GB	Alternative	Good balance of quality and speed. Fallback option.
qwen2:7b	7B	~6GB	Multilingual	Essential for non-English corpora (FR, DE, ES).
deepseek-v2:16b	16B	~12GB	High-Quality	For final report synthesis when quality is paramount.

**The MoE Advantage:**

Mixture-of-Experts models like `granite3.1-moe:3b` are particularly suited for edge compute:

- **Sparse activation:** Only a subset of parameters (experts) activate per token
- **Reduced inference cost:** Effective compute of ~2B despite 3B total parameters
- **Specialization:** Different experts handle different content types

#### Dual-LLM Architecture:



This dual-LLM pattern achieves 61%+ acceptance rate on first pass, with the tutor providing quality assurance without human intervention.

**Recommendation:** Use `granite3.1-moe:3b` as worker for bulk extraction (fast, low VRAM), `mistral:7b-instruct` as tutor for refinement (higher quality validation). Reserve larger models (16B+) for final report synthesis only.

### 3. The KOAS Paradigm: Kernels Compute, LLMs Interpret

#### 3.1 Separation of Concerns

The fundamental architectural principle of KOAS is:

"Kernels compute, LLMs interpret."

This separation provides:

Property	Kernels	LLMs
Determinism	✓ Identical inputs → identical outputs	✗ Stochastic
Auditability	✓ Fully traceable	⚠ Requires caching
Reproducibility	✓ Hash-verifiable	✓ With cache replay
Speed	✓ Fast (compute-bound)	⚠ Slower (inference)
Sovereignty	✓ No external deps	✓ With local deployment

#### 3.2 Kernel Architecture

KOAS kernels are stateless functions with explicit dependencies:

```

@kernel("doc_extract", stage=2, requires=["doc_metadata", "doc_concepts"])
def doc_extract(input: KernelInput) -> KernelOutput:
    """
    Extract key sentences from documents with quality scoring.

    INVARIANT: This kernel performs NO LLM calls.
    All processing is deterministic and reproducible.
    """
    # Pure computation: scoring, filtering, deduplication
    ...

```

### Key properties:

1. **Explicit dependencies** — DAG-based execution order
2. **Input hashing** — Automatic cache invalidation on change
3. **Output schemas** — Structured JSON with validation
4. **No side effects** — Kernels don't modify input data

### 3.3 LLM Boundary

LLM calls are confined to specific synthesis kernels:

Kernel	LLM Role	Isolation Mechanism
doc_func_extract	Extract functionalities from SPD	Cached, replayable
doc_summarize	Generate natural language summaries	Cached, replayable
doc_summarize_tutored	Dual-LLM refinement	Cached, replayable

All other kernels (15+) operate without LLM involvement.

### 3.4 Mathematical Foundation

The KOAS pipeline can be modeled as a directed acyclic graph (DAG)  $G = (V, E)$  where:

- $V$  = set of kernels
- $E$  = dependency edges
- For each kernel  $k \in V$ :  $\text{output}_k = f_k(\text{inputs}_k, \text{config}_k)$

**Reproducibility Theorem:** For a fixed configuration  $C$  and input corpus  $D$ , the final report  $R$  is deterministic if all LLM calls are replaced by cached responses.

$$R = \text{Pipeline}(D, C, \text{LLM\_Cache}) = \text{constant}$$

This enables **verification without LLM**: an auditor can replay the entire pipeline using only cached data.

## 4. KOAS Kernel Inventory and Capabilities

### 4.1 Complete Kernel Registry

RAGIX/KOAS provides approximately **50 specialized kernels** organized into three families. Each kernel is classified by its **determinism level**, which directly impacts auditability and reproducibility.

**Table 1. KOAS Kernel Inventory by Family**

## CODE AUDIT KERNELS ( RAGIX\_KERNELS/AUDIT/ )

Kernel	Stage	Determinism	Purpose	Output
ast_scan	1	✓ Full	Parse source code into AST	Symbols, classes, methods, fields
metrics	1	✓ Full	Compute code metrics	LOC, complexity, maintainability
dependency	1	✓ Full	Build dependency graph	Import/call relationships
services	1	✓ Full	Detect service boundaries	Microservice catalog
coupling	2	✓ Full	Compute coupling metrics	Afferent/efferent coupling, instability
hotspots	2	✓ Full	Identify complexity hotspots	Risk-ranked method list
dead_code	2	✓ Full	Detect unreachable code	Unused symbols, orphan files
entropy	2	✓ Full	Measure code entropy	Information density metrics
partition	2	✓ Full	Graph-based partitioning	Module decomposition proposals
drift	2	⚠ Partial	Detect architectural drift	Pattern violations
timeline	2	✓ Full	Git history analysis	Change frequency, authors
risk	2	✓ Full	Compute risk scores	Per-file risk assessment
volumetry	2	✓ Full	Operational volumetry	Flow weights, peak patterns
module_group	2	✓ Full	Functional grouping	Module-to-flow mapping
risk_matrix	2	✓ Full	Weighted risk matrix	Risk = f(complexity, volumetry)
stats_summary	3	✓ Full	Aggregate statistics	Summary tables
section_executive	3	⚠ Partial	Executive summary	Natural language (LLM)
section_overview	3	✓ Full	Technical overview	Structured sections
section_recommendations	3	⚠ Partial	Recommendations	Prioritized action list
section_drift	3	✓ Full	Drift report section	Violation details
report_assemble	3	✓ Full	Assemble final report	Markdown document

## SECURITY KERNELS ( RAGIX\_KERNELS/SECURITY/ )

Kernel	Stage	Determinism	Purpose	Output
net_discover	1	✓ Full	Network enumeration	Asset inventory
port_scan	1	✓ Full	Service detection	Open ports, services
dns_enum	1	✓ Full	DNS analysis	Subdomains, records
config_parse	1	✓ Full	Parse firewall configs	Iptables, Cisco, pfSense rules
ssl_analysis	2	✓ Full	TLS/certificate audit	Cipher suites, expiry, vulnerabilities
vuln_assess	2	✓ Full	CVE mapping	Vulnerability matches
web_scan	2	✓ Full	Web application scan	OWASP findings
compliance	2	✓ Full	Compliance checking	ANSSI/NIST/CIS mappings
risk_network	2	✓ Full	Network risk scoring	Per-asset risk
section_security	3	⚠ Partial	Security report	Findings narrative

## DOCUMENT KERNELS ( RAGIX\_KERNELS/DOCS/ )

Kernel	Stage	Determinism	Purpose	Output
doc_metadata	1	✓ Full	Document inventory	File list, chunk counts
doc_concepts	1	✓ Full	Concept extraction	Knowledge graph traversal
doc_structure	1	✓ Full	Section detection	Heading hierarchy
doc_quality	1	✓ Full	Quality scoring	MRI, SRI metrics
doc_cluster	2	⚠ Partial	Hierarchical clustering	Document groups
doc_cluster_leiden	2	⚠ Partial	Community detection	Leiden partitions
doc_cluster_reconcile	2	✓ Full	Merge clustering results	Unified grouping
doc_extract	2	✓ Full	Key sentence extraction	Quality-filtered excerpts
doc_coverage	2	✓ Full	Concept coverage analysis	Gaps, overlaps
doc_compare	2	✓ Full	Cross-document comparison	Discrepancies, terminology
doc_func_extract	2	✗ LLM	Functionality extraction	Structured requirements
doc_pyramid	3	✓ Full	Hierarchical summary	4-level pyramid
doc_summarize	3	✗ LLM	Natural language summaries	Per-document summaries
doc_summarize_tutored	3	✗ LLM	Dual-LLM refinement	Quality-gated summaries
doc_visualize	3	✓ Full	Generate visualizations	Word clouds, graphs
doc_report_assemble	3	✓ Full	Assemble report sections	Structured markdown
doc_final_report	3	✓ Full	Final report generation	Complete audit report

## INFRASTRUCTURE KERNELS

Kernel	Determinism	Purpose
preflight	✓ Full	System readiness check
orchestrator	✓ Full	DAG execution engine
cache	✓ Full	LLM and kernel caching
llm_wrapper	⚠ Partial	LLM call boundary

### Legend:

- **✓ Full** — Identical inputs always produce identical outputs
- **⚠ Partial** — Deterministic with fixed seeds or cached responses
- **✗ LLM** — Requires LLM, deterministic only with cache replay

### Seed Locations for Partial-Determinism Kernels:

Kernel	Seed Source	Audit Trail Field	Reproducibility
doc_cluster	config.clustering_seed	kernel_execution[].seed	Same seed → same clusters
doc_cluster_leiden	config.leiden_seed	kernel_execution[].seed	Same seed → same partitions
partition	config.partition_seed	kernel_execution[].seed	Same seed → same modules
section_executive	LLM cache	llm_calls[].call_hash	Same cache → same text
section_recommendations	LLM cache	llm_calls[].call_hash	Same cache → same text
section_security	LLM cache	llm_calls[].call_hash	Same cache → same text
llm_wrapper	LLM cache	llm_calls[].call_hash	Same cache → same response

**To reproduce a run:** Extract seeds from audit trail and pass via `--seed-file=audit_trail.json`.

## 4.2 Determinism Spectrum: From Pure Computation to LLM Synthesis

KOAS kernels span a spectrum of determinism, enabling different audit and reproducibility guarantees.

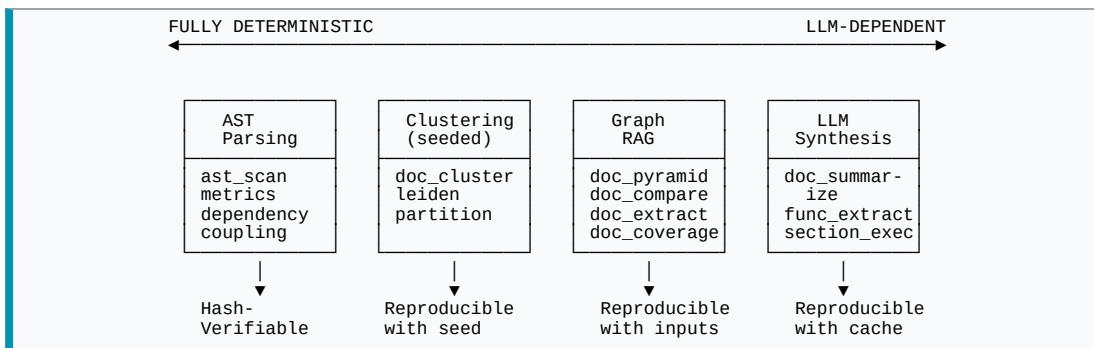


Figure 3. Determinism Spectrum

## 4.3 Capabilities by Analysis Mode

### MODE A: FULLY DETERMINISTIC (AST-BASED CODE AUDIT)

**Use case:** Security audits, compliance verification, technical debt assessment

**Activated kernels:** `ast_scan` → `metrics` → `dependency` → `coupling` → `hotspots` → `dead_code` → `risk` → `report_assemble`

#### Capabilities:

- **Root cause identification:** Trace bugs to specific code locations
- **Complexity hotspots:** Identify methods requiring refactoring
- **Coupling analysis:** Detect architectural violations
- **Dead code detection:** Find unused symbols for removal
- **Dependency mapping:** Full call graph visualization
- **Technical debt estimation:** Remediation effort in hours

**Example output (illustrative values, actual numbers vary by codebase):**

```
{
```

```

    "hotspots": [
      {
        "symbol": "PaymentProcessor.processTransaction",
        "file": "src/payments/processor.py:142",
        "complexity": 47,
        "risk_score": 0.89,
        "recommendation": "Extract validation logic into separate methods"
      }
    ],
    "dead_code": {
      "unused_methods": 23,
      "orphan_files": 4,
      "estimated_loc_removable": 1847
    }
  }
}

```

## MODE B: SEMI-DETERMINISTIC (GRAPH RAG + PYRAMIDAL DECOMPOSITION)

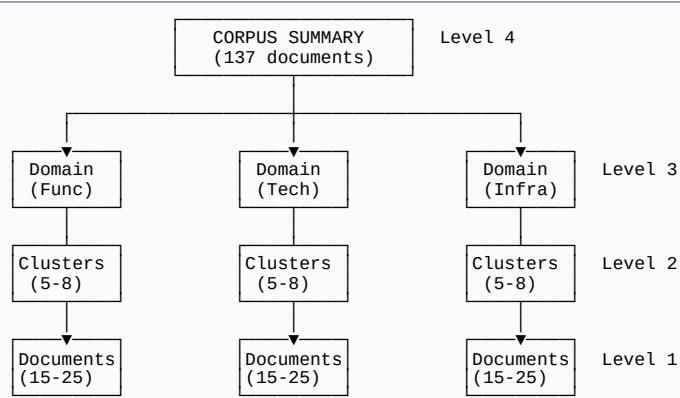
**Use case:** Large corpus analysis, concept extraction, cross-document comparison

**Activated kernels:** doc\_metadata → doc\_concepts → doc\_cluster → doc\_cluster\_leiden → doc\_extract → doc\_compare → doc\_pyramid

**Capabilities:**

- **Concept extraction:** Identify key themes across corpus
- **Analogy and patterns:** Detect similar structures in different documents
- **Repetition detection:** Find redundant content
- **Terminology variations:** Track inconsistent naming
- **Cross-reference validation:** Verify document links
- **Hierarchical summarization:** Document → Cluster → Domain → Corpus

**Pyramidal decomposition:**



**Example output (discrepancies) (illustrative values):**

```

{
  "discrepancies": [
    {
      "type": "terminology_variation",
      "base_term": "authentification",
      "variants": ["authentication", "authent.", "auth."],
      "occurrences": 47
    },
    {
      "type": "missing_reference",
      ...
    }
  ]
}

```

```

        "source": "DOC-042",
        "reference": "SPD-99",
        "context": "See SPD-99 for protocol details"
    }
]
}

```

## MODE C: LLM-ENHANCED (FULL DOCUMENT SYNTHESIS)

**Use case:** Executive summaries, natural language reports, functionality extraction

**Activated kernels:** All of Mode B + `doc_func_extract` → `doc_summarize_tutored` → `doc_final_report`

**Capabilities:**

- **Natural language summaries:** Human-readable document descriptions
- **Functionality extraction:** Structured requirements from specifications
- **Executive synthesis:** High-level insights for decision-makers
- **Quality-gated generation:** Dual-LLM validation

**Example output (functionality) (illustrative values):**

```
{
  "functionalities": [
    {
      "id": "FUNC-016-01",
      "name": "User Authentication via SSO",
      "description": "The system shall authenticate users through corporate SSO...",
      "actors": ["End User", "Administrator"],
      "trigger": "Login attempt",
      "references": ["DOC-034", "DOC-035"],
      "category": "security"
    }
  ]
}
```

## 4.4 Automatic Report Generation with Configurable Isolation

KOAS generates reports through an assembly pipeline where **human control** is paramount. Each stage offers configuration points for isolation level.

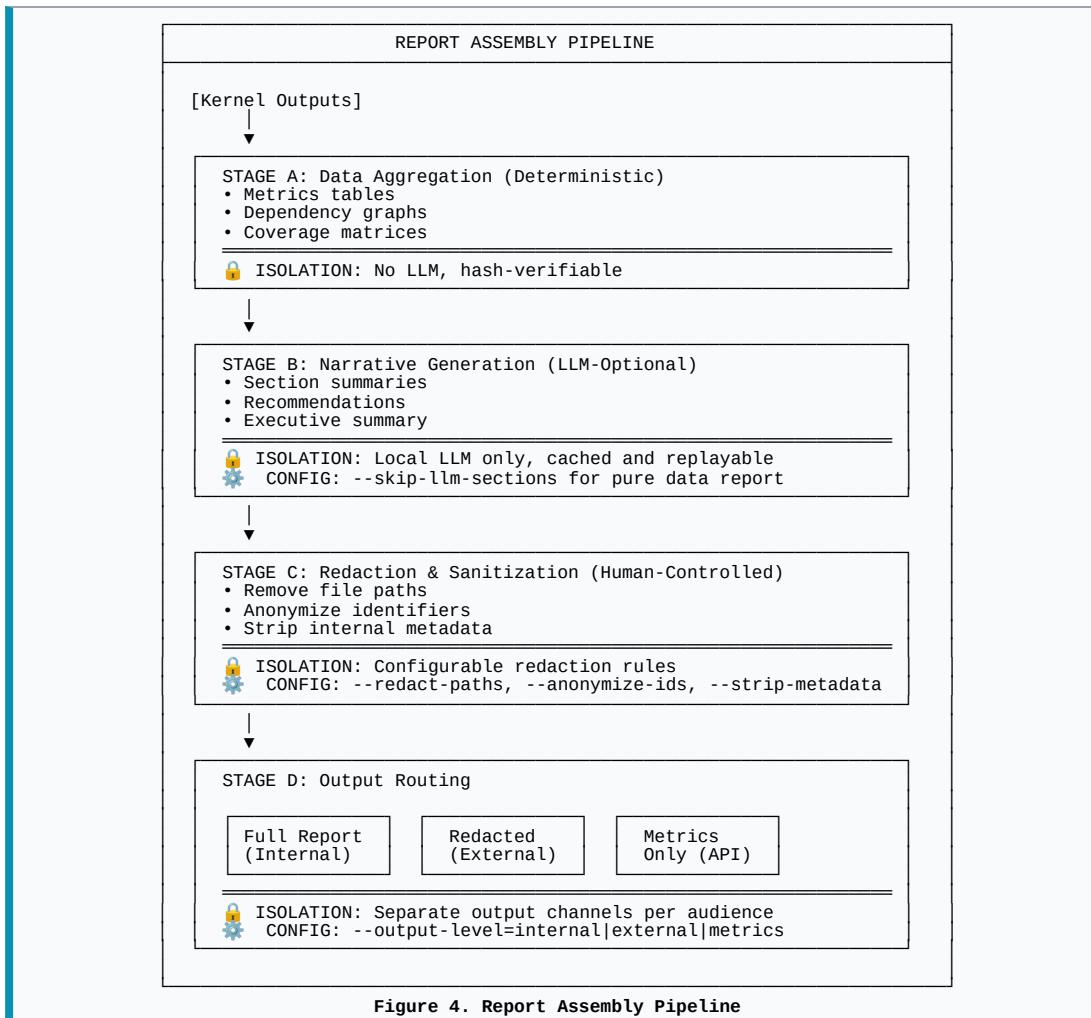


Figure 4. Report Assembly Pipeline

### Isolation Configuration Matrix:

Scenario	LLM Sections	File Paths	Identifiers	Metadata	Output
Internal Audit	✓ Include	✓ Include	✓ Include	✓ Include	Full
External Delivery	✓ Include	✗ Redact	⚠ Anonymize	✗ Strip	Redacted
Orchestrator View	✗ Skip	✗ Redact	✗ Redact	✗ Strip	Metrics
Compliance Archive	✓ Include	✓ Include	✓ Include	✓ Include	Full + Attestation

### 4.5 Human Control Points

Throughout the pipeline, humans retain full control:

Control Point	Options	Default
Kernel selection	Include/exclude specific kernels	All enabled
LLM model choice	Swap worker/tutor models	granite3.1-moe:3b / mistral:7b
Cache policy	write_through / read_only / off	write_through
Redaction rules	Custom regex patterns	None
Output channels	Internal / external / metrics	Internal
Approval gates	Require human sign-off per stage	Disabled
Audit verbosity	Minimal / standard / verbose	Standard

### Example: Maximum Isolation Configuration

```
ragix-koas run \
    --workspace ./audit \
    --all \
    --llm-cache=read_only \
    --kernel-cache=read_only \
    --skip-llm-sections \
    --redact-paths \
    --anonymize-ids \
    --strip-metadata \
    --output-level=metrics \
    --require-approval
```

This produces a metrics-only report suitable for external orchestrators, with all sensitive information removed.

## 5. Policy Enforcement Mechanisms

### 5.1 Forced LLM Caching with Canonicalization

For cache replay to work reliably, LLM requests must be **canonicalized** before hashing.

#### CANONICALIZATION RULES

```
def canonicalize_llm_request(request: dict) -> str:
    """Produce stable JSON for cache key computation."""
    canonical = {
        "model": request["model"],
        "temperature": request.get("temperature", 0.0),
        "template_id": request.get("template_id"),
        "template_version": request.get("template_version"),
        "messages": _canonicalize_messages(request["messages"]),
    }
    # Sorted keys, no whitespace variance
    return json.dumps(canonical, sort_keys=True, separators=(',', ':'))

def _canonicalize_messages(messages: List[dict]) -> List[dict]:
    """Normalize message content."""
    result = []
    for msg in messages:
        content = msg["content"]
        # Normalize whitespace
        content = re.sub(r'\s+', ' ', content).strip()
        # Remove volatile fields (timestamps, run IDs)
        content = re.sub(r'run_\d{8}_\d{6}_[a-f0-9]+', 'RUN_ID', content)
        result.append({"role": msg["role"], "content": content})
    return result
```

#### MERKLE ROOT FOR PYRAMIDAL SYNTHESIS

When summarizing aggregated content (cluster → domain → corpus), children must be ordered deterministically:

```
def compute_inputs_merkle_root(children: List[dict]) -> str:
    """Compute Merkle root of child inputs for cache key."""
    # Step 1: Sort children deterministically
    sorted_children = sorted(children, key=lambda c: (c["file_path"],
        c.get("chunk_index", 0)))

    # Step 2: Hash each child
    child_hashes = [sha256(c["content_hash"]) for c in sorted_children]
```

```

# Step 3: Build Merkle tree
while len(child_hashes) > 1:
    if len(child_hashes) % 2 == 1:
        child_hashes.append(child_hashes[-1]) # Duplicate last if odd
    child_hashes = [sha256(child_hashes[i] + child_hashes[i+1])
                    for i in range(0, len(child_hashes), 2)]

return child_hashes[0] if child_hashes else sha256("")

```

### Cache Key Structure:

```
{
  "call_hash": "sha256(canonical_request)",
  "inputs_merkle_root": "sha256(ordered_children_tree)",
  "template": "docs/domain_summary@1.3",
  "model": "granite3.1-moe:3b",
  "temperature": 0.3
}
```

## 5.2 Cache Modes

The `--llm-cache` flag controls LLM interaction:

Mode	Behavior	Use Case
<code>write_through</code>	Call LLM, cache response	Normal operation
<code>read_only</code>	Use cache only, fail on miss	Verification, air-gapped
<code>read_prefer</code>	Cache first, LLM fallback	Incremental updates
<code>off</code>	No caching	Fresh analysis (rare)

### Air-gapped operation:

```

# Populate cache (requires LLM)
ragix-koas run --workspace ./audit --all --llm-cache=write_through

# Transfer cache to air-gapped system
rsync -av ./audit/.KOAS/cache/ /air-gapped/audit/.KOAS/cache/

# Verify on air-gapped system (no LLM required)
ragix-koas run --workspace /air-gapped/audit --all --llm-cache=read_only

```

## 5.3 Kernel Cache Modes

The `--kernel-cache` flag controls kernel output caching:

Mode	Behavior	Use Case
<code>write_through</code>	Execute kernel, cache output	Normal operation
<code>read_only</code>	Use cached output, skip execution	Ultra-fast replay
<code>off</code>	Always execute	After code changes

### Combined for maximum speed:

```

# Full replay (no computation, no LLM)
ragix-koas run --workspace ./audit --all \
  --llm-cache=read_only \
  --kernel-cache=read_only
# Result: ~45 seconds for 137 documents (vs ~12 minutes full run)

```

## 5.4 Input Validation and Sanitization

Before processing, documents undergo validation. **Critical distinction:**

### BOILERPLATE DETECTION (PURE DOCS MODE ONLY)

**Scope:** Applies ONLY to document corpora (`--type docs`). NEVER applied to code files or mixed mode.

**Purpose:** Filter non-informative content (release notes, infrastructure notation) from representative excerpts.

**Protection:** Fenced code blocks (`````) and inline code (`'`'`) are **excluded** from boilerplate matching:

```
def _apply_boilerplate_filter(text: str, patterns: List[str]) -> str:
    """Apply boilerplate patterns, protecting code blocks."""
    # Step 1: Extract and protect code blocks
    code_blocks = re.findall(r'```[\s\S]*?```', text)
    protected = text
    for i, block in enumerate(code_blocks):
        protected = protected.replace(block, f'__CODE_BLOCK_{i}__')

    # Step 2: Apply boilerplate patterns to non-code content
    for pattern in patterns:
        protected = re.sub(pattern, '', protected, flags=re.IGNORECASE)

    # Step 3: Restore code blocks
    for i, block in enumerate(code_blocks):
        protected = protected.replace(f'__CODE_BLOCK_{i}__', block)

    return protected
```

**Language Packs:** Boilerplate vocabulary is language-specific:

Language	Config Key	Example Patterns
French	<code>boilerplate_vocab_fr</code>	<code>Sommaire, Table des matières, Confidential</code>
English	<code>boilerplate_vocab_en</code>	<code>Table of Contents, Confidential, Draft</code>
German	<code>boilerplate_vocab_de</code>	<code>Inhaltsverzeichnis, Vertraulich</code>
Changelog	<code>boilerplate_changelog</code>	<code>X.X.X.X MERGED, VM1 VM2</code>

Selected via `--language` flag or corpus auto-detection.

### REPRESENTATIVE EXTRACTION QUALITY (DOC\_EXTRACT)

**Scope:** Applies to `doc_extract` kernel in all modes.

**Method:** Boilerplate is penalized via `_score_sentence_quality()`, NOT by cleaning/modifying text:

```
def _score_sentence_quality(sentence: str, config: QualityConfig) -> float:
    """Score sentence quality. Original text is NEVER modified."""
    score = 0.5  # Base score

    # Penalties (reduce score)
    if _matches_boilerplate(sentence, config):
        score -= 0.3  # Heavy penalty for boilerplate
    if _is_truncated(sentence):
        score -= 0.2
```

```

# Bonuses (increase score)
if _has_technical_terms(sentence, config):
    score += 0.15
if _has_action_verbs(sentence):
    score += 0.1
if _optimal_length(sentence):
    score += 0.1

return max(0.0, min(1.0, score))

```

**Key invariant:** The sentence text itself is preserved verbatim. Low-scoring sentences are excluded from representative excerpts, not modified.

## OTHER VALIDATIONS

1. **Metadata stripping** — Separate front-matter from content (see §5.4)
2. **Encoding normalization** — UTF-8 standardization
3. **Path sanitization** — Prevent traversal attacks

### 5.5 Metadata Stripping Contract

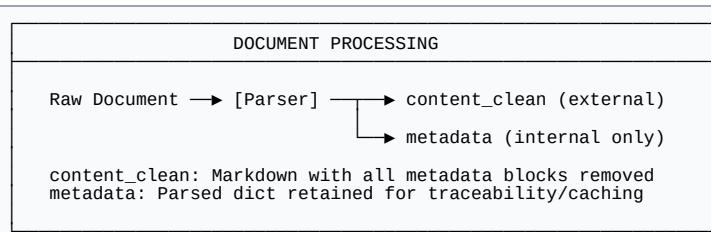
#### Definition of Metadata (Pipeline Invariant)

The following are classified as **internal metadata** and MUST be stripped before external output:

Type	Pattern	Example
YAML front-matter	<code>^---\n...---</code> at file start	<code>title: ... \nllm_trace: ...</code>
TOML front-matter	<code>^+++ \n... \n+++</code> at file start	<code>[meta]\ncall_hash = "..."</code>
JSON front-matter	<code>^```json\n{...}\n``` \n</code> at file start	Embedded config blocks
Fenced metadata	<code>```metadata\n... \n```</code>	Provenance blocks
HTML provenance	<code>&lt;!-- PROVENANCE ... --&gt;</code>	Inline trace markers
Internal keys	Any of: <code>llm_trace</code> , <code>call_hash</code> , <code>inputs_merkle_root</code> , <code>run_id</code> , <code>endpoint</code> , <code>model</code> , <code>cache_status</code>	Audit fields

#### Two-Channel Output Model

Every document processing pipeline produces two representations:



#### Usage Rules:

Context	Allowed Source	Rationale
User-facing reports	content_clean only	No provenance leak
LLM summarization	content_clean only	No meta-confusion
Embedding/retrieval	content_clean only	Clean vectors
Internal audit	Both	Full traceability
Cache key computation	metadata	Provenance tracking

### Build-Time Validation (Denylist Check)

The report assembly pipeline includes a mandatory denylist scan:

```
DENYLIST_KEYS = [
    "llm_trace", "call_hash", "inputs_merkle_root",
    "run_id", "endpoint", "model", "cache_status",
    "prompt_tokens", "completion_tokens", "digest"
]

def validate_external_report(report_path: Path) -> bool:
    """Fail build if denylist keys appear in external report."""
    content = report_path.read_text()
    for key in DENYLIST_KEYS:
        if key in content:
            raise SecurityViolation(f"Denylist key '{key}' found in external report")
    return True
```

This check runs automatically when `--output-level=external` is specified.

### 5.6 Output Isolation

Audit artifacts are written to `.KOAS/` directory:

```
workspace/
└── src/
    └── doc1.pdf
    └── doc2.docx
└── .KOAS/
    ├── cache/
    ├── stage1/
    ├── stage2/
    ├── stage3/
    ├── final_report.md
    └── audit_trail.json
# Input documents (READ-ONLY)
# Output directory (WRITE)
# LLM and kernel caches
# Collection outputs
# Analysis outputs
# Synthesis outputs
# Generated report
# Provenance log
```

This separation ensures:

- Input documents are never modified
- Outputs don't contaminate future RAG indexing
- Clear audit boundary

### 5.7 External vs Internal Output Contract

The `--output-level` flag enforces isolation rules at build time. This is a **hard contract**, not a guideline.

**Table 2. Output Level Contract**

Level	Metadata	File Paths	Identifiers	Excerpts	Enforced By
internal	✓ Include	✓ Include	✓ Include	✓ Include	Default
external	✗ Strip	✗ Redact	⚠ Anonymize	✓ Include	CLI + Build
orchestrator	✗ Strip	✗ Redact	✗ Redact	✗ Metrics only	CLI + Build
compliance	✓ Include	✓ Include	✓ Include	✓ Include	+ Attestation

### CLI Enforcement:

```

# External delivery (client-facing)
ragix-koas run --workspace ./audit --all --output-level=external
# Enforces: strip metadata, redact paths, anonymize IDs
# Build FAILS if denylist keys detected in output

# Orchestrator view (Claude/GPT-4 integration)
ragix-koas run --workspace ./audit --all --output-level=orchestrator
# Enforces: metrics only, no text content
# Returns: kernel names, timing, success/failure, counts

# Compliance archive (full audit trail)
ragix-koas run --workspace ./audit --all --output-level=compliance
# Enforces: full content + sovereignty attestation
# Includes: hash chains, provenance markers, timestamps

```

### Build-Time Validation:

When `--output-level=external`:

1. Denylist scan runs on all generated markdown
2. Path patterns (`/home/`, `/path/to/`, `C:\`) trigger redaction or failure
3. Internal ID formats (`F000XXX`, `run_YYYYMMDD_`) are anonymized

### Example: Anonymization Transform

```

Before: "Document F000142 (src/specs/SPD-16.docx) shows..."
After: "Document [DOC-A] ([PATH-REDACTED]) shows..."

```

**Non-Negotiable:** Absent `--output-level` flag defaults to `internal`. External outputs require explicit declaration.

## 6. Operational Modes

### 6.1 Mode 1: Documents Only

Analyze a corpus of technical documents without code.

**Applicable to:** Specification audits, policy reviews, contract analysis.

```

ragix-koas init --project ./specs --name "Contract Review" --type docs
ragix-koas run --workspace ./specs/.KOAS --all

```

**Kernels activated:** `doc_metadata`, `doc_concepts`, `doc_structure`, `doc_cluster`, `doc_extract`, `doc_pyramid`, `doc_summarize`, `doc_final_report`

## 6.2 Mode 2: Documents + Code

Combined analysis of specifications and implementation.

**Applicable to:** Compliance verification, specification-implementation gap analysis.

```
ragix-koas init --project ./system --name "Compliance Audit" --type mixed
ragix-koas run --workspace ./system/.KOAS --all --include-code
```

**Additional kernels:** `ast_scan`, `metrics_compute`, `coverage_analyze`, `drift_detect`

## 6.3 Mode 3: Code Only

Pure code audit without documentation.

**Applicable to:** Security audits, technical debt assessment, refactoring planning.

```
ragix-koas init --project ./codebase --name "Security Audit" --type code
ragix-koas run --workspace ./codebase/.KOAS --all
```

**Kernels activated:** Standard KOAS code audit pipeline.

## 7. Gray Environment Deployment

### 7.1 Definition

A **gray environment** is an operational context where:

1. Network connectivity exists but is untrusted
2. External API calls are prohibited by policy
3. Data must not leave the local system
4. Auditability is mandatory

This differs from:

- **Green environment:** Full connectivity, cloud services allowed
- **Air-gapped (black):** No network connectivity

### 7.2 Network Configuration

**Firewall rules for gray deployment:**

```
# Allow only localhost for LLM
iptables -A OUTPUT -d 127.0.0.1 -j ACCEPT
iptables -A OUTPUT -d ::1 -j ACCEPT

# Block known LLM API endpoints
iptables -A OUTPUT -d api.openai.com -j DROP
iptables -A OUTPUT -d api.anthropic.com -j DROP
iptables -A OUTPUT -d generativelanguage.googleapis.com -j DROP

# Log any other outbound attempts (for audit)
iptables -A OUTPUT -j LOG --log-prefix "OUTBOUND_ATTEMPT: "
iptables -A OUTPUT -j DROP
```

### 7.3 Verification Procedure

Before processing sensitive data, verify sovereignty:

```
# 1. Check Ollama is local
curl -s http://127.0.0.1:11434/api/tags | jq '.models[].name'

# 2. Verify no external endpoints in config
grep -r "api\." ~/.ragix/ /etc/ragix/ 2>/dev/null | grep -v localhost

# 3. Test network isolation
timeout 5 curl -s https://api.openai.com/v1/models && echo "FAIL: External
reachable" || echo "OK: External blocked"

# 4. Run RAGIX system check
ragix-koas status --sovereignty-check
```

### 7.4 Sovereignty Attestation

The audit trail includes sovereignty attestation:

```
{
  "sovereignty": {
    "hostname": "audit-workstation-01",
    "user": "auditor",
    "platform": "Linux 6.8.0-90-generic",
    "python_version": "3.12.12",
    "llm_endpoint": "http://127.0.0.1:11434",
    "llm_local": true,
    "models_used": [
      {"name": "granite3.1-moe:3b", "digest": "b43d80d7fca7", "role": "worker"},
      {"name": "mistral:7b-instruct", "digest": "6577803aa9a0", "role": "tutor"}
    ],
    "external_calls": 0,
    "attestation": "All processing performed locally. No data sent to external
services."
  }
}
```

## 8. Attestation and Audit Trail

### 8.1 Audit Trail Structure

Every KOAS run produces a comprehensive audit trail:

```
{
  "_meta": {
    "version": "1.0.0",
    "generated_at": "2026-01-29T14:05:10Z",
    "generator": "KOAS Document Summarization v0.64.2"
  },
  "run_id": "run_20260129_140510_8a3f2c",
  "sovereignty": { ... },
  "configuration": {
    "project_root": "/path/to/audit",
    "language": "fr",
    "llm_model": "granite3.1-moe:3b",
    "llm_cache_mode": "write_through",
    "kernel_cache_mode": "write_through"
  },
}
```

```

"kernel_execution": [
    {
        "name": "doc_metadata",
        "stage": 1,
        "success": true,
        "execution_time_s": 0.01,
        "input_hash": "sha256:6c9d206bc3aa10d3...",
        "output_hash": "sha256:4b35b47917038e35...",
        "llm_calls": 0
    },
    // ... all kernels
],
"llm_calls": [
    {
        "call_hash": "sha256:a3f2c8d1...",
        "kernel": "doc_summarize",
        "model": "granite3.1-moe:3b",
        "prompt_tokens": 1247,
        "completion_tokens": 312,
        "cache_status": "miss",
        "timestamp": "2026-01-29T14:06:23Z"
    }
],
"checksums": {
    "input_corpus": "sha256:...",
    "final_report": "sha256:...",
    "audit_trail": "sha256:..."
}
}
}

```

## 8.2 Hash Chain Integrity

Logs are protected by SHA-256 hash chaining:

```

Entry[0]: { data, hash_0 = SHA256(data) }
Entry[1]: { data, hash_1 = SHA256(data || hash_0) }
Entry[n]: { data, hash_n = SHA256(data || hash_{n-1}) }

```

### Verification:

```

ragix verify --audit-trail ./audit/.KOAS/audit_trail.json
# Output: "Chain integrity verified. 247 entries, no tampering detected."

```

## 8.3 Provenance Tracking

Every generated statement can be traced:

```

## Domain: Infrastructure Management

This domain covers network configuration and system monitoring...

<!-- PROVENANCE
call_hash: sha256:a3f2c8d1e5b7...
input_documents: [F000023, F000024, F000031]
kernel: doc_summarize
timestamp: 2026-01-29T14:06:23Z
-->

```

**Metadata stripping:** Provenance markers are stripped from user-facing outputs but preserved in internal audit artifacts.

## 9. Practical Implementation

### 9.1 Prerequisites

#### Hardware:

- CPU: 8+ cores recommended
- RAM: 16GB minimum, 32GB recommended
- Storage: SSD, 50GB+ free space
- GPU: Optional but accelerates inference (NVIDIA 8GB+ VRAM)

#### Software:

```
# Ollama installation
curl -fsSL https://ollama.com/install.sh | sh

# Model download (offline-capable)
ollama pull granite3.1-moe:3b
ollama pull mistral:7b-instruct

# RAGIX installation
pip install ragix[full]
```

### 9.2 Initialization

```
# Create audit workspace
ragix-koas init \
  --project /path/to/documents \
  --name "Technical Audit 2026" \
  --language fr \
  --type docs

# Verify configuration
cat /path/to/documents/.KOAS/manifest.yaml
```

### 9.3 Execution

```
# Full pipeline (populates caches)
ragix-koas run \
  --workspace /path/to/documents/.KOAS \
  --all \
  --llm-cache=write_through \
  --kernel-cache=write_through

# Monitor progress
tail -f /path/to/documents/.KOAS/logs/koas.log
```

### 9.4 Verification

```
# Replay without LLM (verification mode)
ragix-koas run \
  --workspace /path/to/documents/.KOAS \
  --all \
  --llm-cache=read_only \
  --kernel-cache=read_only

# Verify audit trail
ragix verify --audit-trail /path/to/documents/.KOAS/audit_trail.json

# Generate sovereignty report
ragix-koas report --sovereignty --format pdf
```

## 10. Case Study: Metropolitan Infrastructure Audit

### 10.1 Context

- Client:** Major metropolitan authority
- Scope:** Technical specification corpus for urban infrastructure system
- Documents:** 137 files (DOCX, PDF, PPTX, XLSX)
- Chunks:** 11,809 indexed segments
- Constraint:** Full data sovereignty required

### 10.2 Execution Metrics

**Table 3. Execution Metrics** (example run, actual values vary by corpus size and hardware):

Stage	Duration	Kernels	LLM Calls	Cache Hits
Stage 1	54s	6	0	N/A
Stage 2	~11 min	8	340	39%
Stage 3	42s	6	137	61%
<b>Total</b>	<b>~13 min</b>	<b>20</b>	<b>477</b>	<b>52%</b>

**Replay (cached):** ~45 seconds (example)

### 10.3 Quality Improvements (v0.64.2)

Metric	Before	After	Improvement
Boilerplate sentences	649	644	-5 (0.8%)
Changelog data in excerpts	Present	Filtered	✓
Infrastructure notation	Present	Filtered	✓
Report path	Wrong	Correct	✓

### 10.4 Sovereignty Verification

```
$ ragix-koas status --sovereignty-check
✓ LLM endpoint: localhost:11434 (local)
✓ External calls: 0
✓ Models verified: granite3.1-moe:3b, mistral:7b-instruct
✓ Cache integrity: 477 entries, chain valid
✓ Attestation: SOVEREIGN
```

## 11. Limitations and Mitigations

### 11.1 Model Quality vs. Cloud Services

**Limitation:** Local 3B-7B models produce lower quality outputs than GPT-4 or Claude Opus.

**Mitigations:**

- Dual-LLM architecture** — Worker + Tutor refinement
- Pre-structured prompts** — Kernels reduce LLM cognitive load
- Iterative refinement** — Multiple passes with quality gates
- Human review** — Final outputs require expert validation

## 11.2 Computational Resources

**Limitation:** Local inference requires significant hardware.

**Mitigations:**

1. **MoE models** — granite3.1-moe:3b activates only 2B params
2. **Kernel caching** — Avoid redundant computation
3. **Batch processing** — Optimize GPU utilization
4. **Incremental updates** — Process only changed documents

## 11.3 Model Updates

**Limitation:** Local models don't receive automatic updates.

**Mitigations:**

1. **Version pinning** — Lock model digests in manifest
2. **Controlled updates** — Manual review before deployment
3. **Regression testing** — Verify outputs against baselines

---

## 12. Conclusion

---

The RAGIX/KOAS framework demonstrates that sovereign LLM operations are achievable for document analysis tasks requiring confidentiality. The key architectural decisions enabling this are:

1. **Local LLM deployment** — Ollama eliminates external API dependencies
2. **Kernel-LLM separation** — Deterministic computation isolated from stochastic inference
3. **Comprehensive caching** — Enables LLM-free replay and verification
4. **Audit trail integrity** — Hash-chained provenance for compliance
5. **Policy enforcement** — Cache modes control data flow

For organizations operating in regulated environments, this approach provides:

- **Compliance** — Data never leaves controlled perimeter
- **Auditability** — Every transformation is traceable
- **Reproducibility** — Results can be independently verified
- **Cost control** — No per-token API charges

The v0.64.2 release demonstrates continuous improvement in content quality while maintaining sovereignty guarantees, with measurable metrics tracking regressions and improvements across versions.

---

## 13. References

---

### Academic Foundations

1. Lewis, P. et al. (2020). "Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks." *NeurIPS 2020*.
2. Edge, D. et al. (2024). "From Local to Global: A Graph RAG Approach to Query-Focused Summarization." *arXiv:2404.16130*.
3. Vitrac, O. (2025). "Virtual/Hybrid R&D Laboratories built with Augmented-AI Agents." *Generative Simulation Initiative*.

## Standards and Regulations

4. **ANSSI** (2024). "Guide d'hygiène informatique." Agence nationale de la sécurité des systèmes d'information.
5. **GDPR** (2016). "General Data Protection Regulation." European Union Regulation 2016/679.
6. **ISO/IEC 27001:2022** — Information security management systems.

## Technical Documentation

7. Ollama Documentation. <https://ollama.com/docs>
8. RAGIX Project Documentation. Internal technical specifications.
9. KOAS Kernel Specifications. [docs/KOAS\\_DOCS.md](#)

## Appendix A: Sovereignty Checklist

Before processing sensitive data, verify:

- Ollama running on localhost:11434
- Required models downloaded and verified
- Firewall rules blocking external LLM APIs
- RAGIX configured with `llm_endpoint: http://127.0.0.1:11434`
- Cache directory on local encrypted storage
- Audit trail enabled
- No cloud backup of `.KOAS/` directory

## Appendix B: Glossary

Term	Definition
<b>Air-gapped</b>	System with no network connectivity
<b>Gray environment</b>	Networked but untrusted, external calls prohibited
<b>KOAS</b>	Kernel-Orchestrated Audit System
<b>MoE</b>	Mixture of Experts (sparse model architecture)
<b>Ollama</b>	Local LLM runtime
<b>RAGIX</b>	Retrieval-Augmented Generative Interactive eXecution
<b>Sovereignty</b>	Data processing within controlled perimeter

RAGIX/KOAS — Adservio Innovation Lab / 2026